

**Department of Computer Science and Engineering (Data Science)**

**Advanced Computational Linguistics**

**Experiment No. 7**

**Name: Umang Kirit Lodaya**

**SAP ID: 60009200032**

**Batch: D11**

**Aim: To study perform various Natural Language Processing task like Text summarization, Sentiment Analysis using available libraries like Hugging Face's Transformers, TensorFlow or PyTorch**

**Introduction:**

**Text Summarization using BERT**

**Extractive Summarization:** BERT can be utilized for extractive summarization, where key sentences or phrases are selected from the original text to form a summary.

Sentence Embedding's: Use BERT to generate embedding's for each sentence in the text.

Similarity Measures: Calculate similarity scores between sentences (e.g., using cosine similarity).

Select Top Sentences: Choose sentences with the highest similarity scores to form the summary.

**Abstractive Summarization:** This involves generating a summary that might not directly include sentences from the original text. BERT can aid in this by fine-tuning a model specifically for abstractive summarization.

**Fine-tuning:** Fine-tune the pre-trained BERT model on a summarization dataset (e.g., CNN/Daily Mail dataset).

**Sequence-to-sequence Model:** Employ techniques like seq2seq models or transformers to generate summaries.

## **Sentiment Analysis using BERT**

**Fine-tuning BERT:** The pre-trained BERT model can be fine-tuned on a sentiment analysis dataset. The model's classification layers can be adjusted for sentiment analysis.

**Dataset Preparation:** Obtain a labeled dataset for sentiment analysis (e.g., IMDB movie reviews, Twitter sentiment dataset).

**Tokenization and Fine-tuning:** Tokenize the text, prepare input sequences, and fine-tune BERT on the sentiment classification task.

**Prediction:** After fine-tuning, use the trained model to predict sentiment labels (positive, negative, neutral) for new text inputs.

### **Tools and Libraries:**

**Hugging Face's Transformers:** This library provides easy access to pre-trained models like BERT and various other NLP-related functionalities for tasks like tokenization, model loading, and fine-tuning.

**TensorFlow or PyTorch:** Use these deep learning frameworks to implement BERT-based models and fine-tuning for specific tasks.

### **Lab Experiment to be performed in this session:**

1. Perform Text Summarization using BERT (BERT summarizer library can be directly installed in python using the following commands `python pip install bert-extractive-summarizer` for the ease of the implementation.)
2. Perform Sentiment Analysis using BERT.

**NAME: UMANG KIRIT LODAYA**

**SAP ID: 60009200032**

**BATCH: D11**

```
In [1]: !pip install -qq transformers
```

```
In [2]: import transformers
        from transformers import BertModel, BertTokenizer, AdamW, get_
        linear_schedule_with_warmup
        import torch

        import numpy as np
        import pandas as pd
        import seaborn as sns
        import matplotlib.pyplot as plt
        from matplotlib import rc
        from sklearn.model_selection import train_test_split
        from sklearn.metrics import confusion_matrix, classification_r
        eport
        from collections import defaultdict
        from textwrap import wrap

        from torch import nn, optim
        from torch.utils.data import Dataset, DataLoader
        import torch.nn.functional as F
        import warnings
        warnings.filterwarnings("ignore")
        %matplotlib inline
```


```
In [3]: device = torch.device("cuda:0" if torch.cuda.is_available()
        se "cpu")
```

```
In [4]: df = pd.read_csv("../input/google-play-storeappreviews/reviews.csv")
df.head()
```

Out[4]:

	reviewId	userName
0	gp:AOqpTOG-nGuDe0w6if400pTvNjNk3MlfnEcJLIB84aR...	Ben Maybe lh.googleusei
1	gp:AOqpTOH85sc18Ajgcgj6-IGmA7Gp34fVsrbyBJ274IZ...	Anthony Duarte lh.googleusei
2	gp:AOqpTOGxyMqOStnhbQ_mLfnLUfd1DHAt5uRXqDNArML...	Through Genesis lh.googleusei
3	gp:AOqpTOE3fQQpkWOMbSHW-DwukAnILBbMwBHEcbReiit...	Sydney Stoll lh.googleusei
4	gp:AOqpTOG2AXuKqrR8FIN43a5BGZY4Iha5SFQZZ6o9vOR...	ChuCannon lh.googleusei

```
In [5]: sns.countplot(df.score)
plt.xlabel('review score');
```



```
In [6]: def sentiment(rating):
        if rating < 2:
            return 0

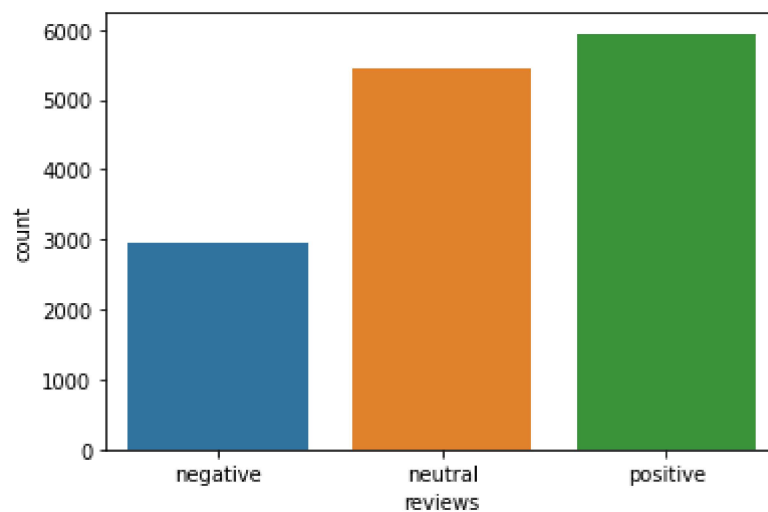
        if rating==3:
            return 1

        if rating > 3:
            return 2

        df['sentiment'] = df.score.apply(sentiment)
```

```
In [7]: ax = sns.countplot(df.sentiment)
        plt.xlabel('reviews')
        ax.set_xticklabels(['negative', 'neutral', 'positive'])
```

```
Out[7]: [Text(0, 0, 'negative'), Text(1, 0, 'neutral'), Text(2, 0, 'positive')]
```



## Tokenizing and Data Preprocessing

```
In [8]: PRE_TRAINED_MODEL_NAME = 'bert-base-cased'
```

```
In [9]: tokenizer = BertTokenizer.from_pretrained(PRE_TRAINED_MODEL_ME)
```

```
Downloading: 29.0/29.0 [00:00<00:00, 100% 1.02kB/s]
```

```
Downloading: 208k/208k [00:00<00:00, 100% 511kB/s]
```

```
Downloading: 426k/426k [00:00<00:00, 100% 717kB/s]
```

```
Downloading: 570/570 [00:00<00:00, 100% 20.8kB/s]
```

**one way**

```
In [10]: # sample_txt = 'When was I last outside? I am stuck at home ;  
r 2 weeks.'  
# tokens = tokenizer.tokenize(sample_txt)  
# token_ids = tokenizer.convert_tokens_to_ids  
# print(f'Sentence: {sample_txt}\n')  
# print(f'Tokens: {tokens}\n')  
# print(f'Token IDs: {token_ids}')
```

[SEP] token has to be inserted at the end of a single input. When a task requires more than one input such as NLI and Q-A tasks, [SEP] token helps the model to understand the end of one input and the start of another input in the same sequence input.

```
In [11]: tokenizer.sep_token, tokenizer.sep_token_id
```

```
Out[11]: ('[SEP]', 102)
```

[CLS] - we must add this token to the start of each sentence, so BERT knows we're doing classification

```
In [12]: tokenizer.cls_token, tokenizer.cls_token_id
```

```
Out[12]: ('[CLS]', 101)
```

The BERT model receives a fixed length of sentence as input. Usually the maximum length of a sentence depends on the data we are working on. For sentences that are shorter than this maximum length, we will have to add paddings (empty tokens) to the sentences to make up the length. In the original implementation, the token [PAD] is used to represent paddings to the sentence.

```
In [13]: tokenizer.pad_token, tokenizer.pad_token_id
```

```
Out[13]: ('[PAD]', 0)
```


### *other way*

```
In [14]: token_lens = []

for text in df.content:
    tokens = tokenizer.encode(text,max_length=128)
    token_lens.append(len(tokens))
```

Truncation was not explicitly activated but 'max\_length' is provided a specific value, please use 'truncation=True' to explicitly truncate examples to max length. Defaulting to 'longest-first' truncation strategy. If you encode pairs of sequences (GLUE-style) with the tokenizer you can select this strategy more precisely by providing a specific strategy to 'truncation'.

```
In [15]: sns.distplot(token_lens)
plt.xlim([0, 200]);
plt.xlabel('Token count');
```



### *and the other way*

```
In [16]: sample_txt = 'When was I last outside? I am stuck at home for  
2 weeks.'
```

```
In [17]: encoding = tokenizer.encode_plus(
        sample_txt,
        max_length = 32,
        add_special_tokens = True, # [CLS] and [SEP]
        pad_to_max_length = True,
        return_token_type_ids = False,
        return_attention_mask = True,
        return_tensors = 'pt' # pt for pytorch
    )
    encoding.keys()
```

```
Out[17]: dict_keys(['input_ids', 'attention_mask'])
```

```
In [18]: # every input is now of size 32 and padded with 0
        print(len(encoding['input_ids'][0]))
        encoding['input_ids'][0]
```

```
32
```

```
Out[18]: tensor([ 101, 1332, 1108,  146, 1314, 1796,  136,  146, 1821,
        5342, 1120, 1313,
        1111,  123, 2277,  119,  102,   0,   0,   0,   0,
        0,   0,   0,   0,   0,   0,   0,   0,   0,   0])
```

```
In [19]: encoding['attention_mask']
```

```
Out[19]: tensor([[1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
        0, 0, 0, 0, 0, 0,
        0, 0, 0, 0, 0, 0, 0, 0]])
```

## Creating Dataset



```
In [20]: class CustomDataset(Dataset):
def __init__(self, reviews, targets, tokenizer, max_len):
    super().__init__()
    self.reviews = reviews
    self.targets = targets
    self.tokenizer = tokenizer
    self.max_len = max_len

def __len__(self):
    return len(self.reviews)

def __getitem__(self, item):
    review = self.reviews[item]
    target = self.targets[item]

    encoding = tokenizer.encode_plus(
        review,
        max_length = self.max_len,
        add_special_tokens = True,
        pad_to_max_length = True,
        return_token_type_ids = False,
        return_attention_mask = True,
        return_tensors = 'pt'
    )

    return {
        'review': review,
        'target': torch.tensor(target, dtype=torch.long),
        'input_ids': encoding['input_ids'].flatten(),
        'attention_mask': encoding['attention_mask'].flatten()
    }
```

```
In [21]: df.shape
```

```
Out[21]: (17082, 13)
```

```
In [22]: df = df[df['sentiment'].notna()]
```

```
In [23]: df.isnull().sum()
```

```
Out[23]: reviewId      0
userName      0
userImage     0
content       0
score         0
thumbsUpCount 0
reviewCreatedVersion 2179
at            0
replyContent  7059
repliedAt     7059
sortOrder     0
appId         0
sentiment     0
dtype: int64
```

```
In [24]: MAX_LEN = 160
        BATCH_SIZE = 16
        RANDOM_SEED = 2002
        EPOCHS = 10
```

```
In [25]: df_train, df_test = train_test_split(df, test_size=0.2, rand
        _state=RANDOM_SEED)
        df_val, df_test = train_test_split(df_test, test_size=0.5, ran
        dom_state=RANDOM_SEED)
```

```
In [26]: df_train.shape, df_val.shape, df_test.shape
```

```
Out[26]: ((11491, 13), (1436, 13), (1437, 13))
```

```
In [27]: ds = CustomDataset(
        reviews=df.content.to_numpy(),
        targets=df.sentiment.to_numpy(),
        tokenizer=tokenizer,
        max_len=MAX_LEN
    )
```

## DataLoading

```
In [28]: def Data_Loader(df, tokenizer, max_len, batch_size):
        ds = CustomDataset(
            reviews=df.content.to_numpy(),
            targets=df.sentiment.to_numpy(),
            tokenizer=tokenizer,
            max_len=max_len
        )

        return DataLoader(
            ds,
            batch_size=batch_size,
            num_workers=4
        )
```

```
In [29]: train_data_loader = Data_Loader(df_train, tokenizer, MAX_LEN
        BATCH_SIZE)
        val_data_loader = Data_Loader(df_val, tokenizer, MAX_LEN, BATC
        H_SIZE)
        test_data_loader = Data_Loader(df_test, tokenizer, MAX_LEN, BA
        TCH_SIZE)
```

```
In [30]: data = next(iter(train_data_loader))
        print(data.keys())

        dict_keys(['review', 'target', 'input_ids', 'attention_mask'])
```



```
In [34]: last_hidden_state.shape # 1 example , 32 elements after padding and 768 hidden units
```

```
Out[34]: torch.Size([1, 32, 768])
```

```
In [35]: pooled_output.shape # pooling procedure on 32 elements
```

```
Out[35]: torch.Size([1, 768])
```

```
In [36]: bert_model.config.hidden_size
```

```
Out[36]: 768
```

## Building Sentiment Classifier

```
In [37]: class SentimentClassifier(nn.Module):
    def __init__(self, n_classes):
        super().__init__()
        self.bert = BertModel.from_pretrained(PRE_TRAINED_MODEL_NAME)
        self.dropout = nn.Dropout(p=0.2)
        self.linear = nn.Linear(self.bert.config.hidden_size, n_classes)
        self.softmax = nn.Softmax(dim=1)

    def forward(self, input_ids, attention_mask):
        _, pooled_output = self.bert(
            input_ids = input_ids,
            attention_mask = attention_mask,
            return_dict = False
        )
        output = self.dropout(pooled_output)
        output = self.linear(output)
        return self.softmax(output)
```

```
In [38]: model = SentimentClassifier(n_classes = 3)
model = model.to(device)
```

Some weights of the model checkpoint at bert-base-cased were not used when initializing BertModel: ['cls.predictions.transform.LayerNorm.bias', 'cls.predictions.decoder.weight', 'cls.predictions.transform.dense.weight', 'cls.predictions.bias', 'cls.predictions.transform.dense.bias', 'cls.predictions.transform.LayerNorm.weight', 'cls.seq\_relationship.weight', 'cls.seq\_relationship.bias']

– This IS expected if you are initializing BertModel from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).

– This IS NOT expected if you are initializing BertModel from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

```

In [39]: optimizer = AdamW(model.parameters(), lr=2e-5, correct_bias=False)

total_steps = len(train_data_loader) * EPOCHS

scheduler = get_linear_schedule_with_warmup(
    optimizer,
    num_warmup_steps = 0,
    num_training_steps = total_steps
)

loss_fn = nn.CrossEntropyLoss().to(device)

```

```

In [40]: def train_model(model, data_loader, loss_fn, optimizer, device, scheduler, n_examples):
    model = model.train()
    losses = []
    correct_preds = 0

    for data in data_loader:
        input_ids = data['input_ids'].to(device)
        attention_mask = data['attention_mask'].to(device)
        targets = data['target'].to(device)

        outputs = model(input_ids, attention_mask)
        _, preds = torch.max(outputs, dim=1)
        loss = loss_fn(outputs, targets)
        correct_preds += torch.sum(preds==targets)
        losses.append(loss.item())

        loss.backward()
        nn.utils.clip_grad_norm(model.parameters(), max_norm=1.0)
    # We're avoiding exploding gradients by clipping the gradients of the model using clip_grad_norm.
    optimizer.step()
    scheduler.step()
    optimizer.zero_grad()

    return correct_preds.double() / n_examples, np.mean(losses)

```

```
In [41]: def eval_model(model, data_loader, loss_fn, device, n_examples):
model = model.eval()
losses = []
correct_preds = 0

with torch.no_grad():
    for data in data_loader:
        input_ids = data['input_ids'].to(device)
        attention_mask = data['attention_mask'].to(device)
        targets = data['target'].to(device)

        outputs = model(input_ids, attention_mask)
        _, preds = torch.max(outputs, dim=1)
        loss = loss_fn(outputs, targets)
        correct_preds += torch.sum(preds==targets)
        losses.append(loss.item())

    return correct_preds.double() / n_examples, np.mean(losses)
```

```
In [42]: %%time
history = defaultdict(list)
best_accuracy = 0

for epoch in range(EPOCHS):
    print(f'Epoch {epoch + 1}/{EPOCHS}')
    print('-' * 10)

    train_acc, train_loss = train_model(model, train_data_loader,
    loss_fn, optimizer, device, scheduler, len(df_train))
    print(f'Train loss {train_loss} accuracy {train_acc}')

    val_acc, val_loss = eval_model(model, val_data_loader, loss_fn,
    device, len(df_val))
    print(f'Val loss {val_loss} accuracy {val_acc}')
    print('\n')

    history['train_acc'].append(train_acc)
    history['train_loss'].append(train_loss)
    history['val_acc'].append(val_acc)
    history['val_loss'].append(val_loss)

    if val_acc > best_accuracy:
#         torch.save(model.state_dict(), 'best_model_state.bin')
        best_accuracy = val_acc
```

Epoch 1/10

---

Train loss 0.8895378269472772 accuracy 0.6472891828387434

Val loss 0.8509273495939043 accuracy 0.6908077994428969

Epoch 2/10

---

Train loss 0.7862448751180328 accuracy 0.7575493864763728

Val loss 0.816420171658198 accuracy 0.7325905292479109

Epoch 3/10

---

Train loss 0.7348239190555249 accuracy 0.81341919763293

Val loss 0.8232506129476759 accuracy 0.7249303621169917

Epoch 4/10

---

Train loss 0.7103045571994383 accuracy 0.8397876599077538

Val loss 0.808188118537267 accuracy 0.7437325905292479

Epoch 5/10

---

Train loss 0.6937786942562905 accuracy 0.8571055608737272

Val loss 0.7955854958958096 accuracy 0.7534818941504178

Epoch 6/10



---

Train loss 0.6844588421979435 accuracy 0.8660690975546079  
Val loss 0.7986241181691488 accuracy 0.7534818941504178

Epoch 7/10

---

Train loss 0.6769924323814137 accuracy 0.8735532155600034  
Val loss 0.7991188453303443 accuracy 0.7506963788300836

Epoch 8/10

---

Train loss 0.670027951786349 accuracy 0.8809503089374292  
Val loss 0.788244284523858 accuracy 0.7618384401114207

Epoch 9/10

---

Train loss 0.6658583255238593 accuracy 0.8852145157079453  
Val loss 0.7841546257336934 accuracy 0.766016713091922

Epoch 10/10

---

Train loss 0.6618157570484782 accuracy 0.8893046732225219  
Val loss 0.784818661875195 accuracy 0.7646239554317549

CPU times: user 34min 41s, sys: 8.72 s, total: 34min 50s


Wall time: 35min 25s

***Increase number of epochs to get better results***

```
In [43]: train_acc = [acc.cpu() for acc in history['train_acc']]  
val_acc = [acc.cpu() for acc in history['val_acc']]
```

```
In [44]: plt.figure(figsize = (8,6))  
plt.plot(train_acc, label='train accuracy')  
plt.plot(val_acc, label='validation accuracy')  
  
plt.title('Training history')  
plt.ylabel('Accuracy')  
plt.xlabel('Epoch')  
plt.legend()  
plt.ylim([0, 1])
```

Out[44]: (0.0, 1.0)



```
In [45]: def get_preds(model, data_loader):
        model = model.eval()

        review_texts = []
        predictions = []
        prediction_probs = []
        real_values = []

        with torch.no_grad():
            for data in data_loader:
                reviews = data['review']
                input_ids = data['input_ids'].to(device)
                attention_mask = data['attention_mask'].to(device)
                targets = data['target'].to(device)

                review_texts.extend(reviews)
                real_values.extend(targets)

                outputs = model(input_ids, attention_mask)
                _, preds = torch.max(outputs, dim=1)

                prediction_probs.extend(outputs)
                predictions.extend(preds)

            predictions = torch.stack(predictions).cpu()
            prediction_probs = torch.stack(prediction_probs).cpu()
            real_values = torch.stack(real_values).cpu()

        return review_texts, predictions, prediction_probs, real_v
        alues
```

```
In [46]: review_texts, predictions, prediction_probs, real_values = g
        _preds(model, test_data_loader)
```

```
In [47]: class_names = ['negative', 'neutral', 'positive']
```

```
In [48]: for i in range(3):
          print('review_texts: ',review_texts[i])
          print('predictions: ',class_names[predictions[i]])
          print('prediction_probs: ',prediction_probs[i])
          print('real_values: ',class_names[real_values[i]])
          print('\n')
```

review\_texts: Not bad

predictions: neutral

prediction\_probs: tensor([2.4137e-05, 9.9993e-01, 4.2055e-05])

real\_values: neutral

review\_texts: Impossible for custom skills to influence multiple characteristics

predictions: neutral

prediction\_probs: tensor([2.2660e-05, 9.9996e-01, 2.1047e-05])

real\_values: neutral

review\_texts: I just cannot be notified. Please. Fix it. So I could give 5 stars.

predictions: negative

prediction\_probs: tensor([9.9994e-01, 1.9388e-05, 3.9770e-05])


real\_values: negative

```
In [49]: print(classification_report(real_values, predictions, target,
                                     names=['negative', 'neutral', 'positive']))
```

	precision	recall	f1-score	support
negative	0.79	0.72	0.75	299
neutral	0.73	0.74	0.73	559
positive	0.81	0.84	0.82	579
accuracy			0.77	1437
macro avg	0.78	0.76	0.77	1437
weighted avg	0.77	0.77	0.77	1437

```
In [50]: def show_confusion_matrix(confusion_matrix):
plt.figure(figsize = (12,10))
hmap = sns.heatmap(confusion_matrix, annot=True, fmt="d",
cmap="Greens")
hmap.yaxis.set_ticklabels(hmap.yaxis.get_ticklabels(), rot
ation=0, ha='right')
hmap.xaxis.set_ticklabels(hmap.xaxis.get_ticklabels(), rot
ation=30, ha='right')
plt.ylabel('True sentiment')
plt.xlabel('Predicted sentiment');

cm = confusion_matrix(real_values, predictions)
df_cm = pd.DataFrame(cm, index=class_names, columns=class_name
s)
show_confusion_matrix(df_cm)
```



## Predicting new text sample

```
In [51]: review_text = 'This notebook is very helpful.'
```

```
In [52]: encoded_review = tokenizer.encode_plus(
    review_text,
    max_length=MAX_LEN,
    add_special_tokens=True,
    return_token_type_ids=False,
    pad_to_max_length=True,
    return_attention_mask=True,
    return_tensors='pt',
)
```

```
In [53]: input_ids = encoded_review['input_ids'].to(device)
         attention_mask = encoded_review['attention_mask'].to(device)

         output = model(input_ids, attention_mask)
         _, prediction = torch.max(output, dim=1)

         print(f'Review text: {review_text}')
         print(f'Sentiment : {class_names[prediction]}')
```

Review text: This notebook is very helpful.

Sentiment : positive